

Lab2 EEGNet

0856029 Yu-Shao Liu

April 12, 2020


1 Introduction

This lab is aim at building a light-weighted CNN model called [EEGNet](#), which is designed for classification problem on EEG. EEGNet is inspired by Xception and uses depthwise separable convolution instead of 3x3 convolution filter to capture features and to decrease parameters. EEGNet is still the state-of-the-art in EEG classification problem. We have presented this work at CECNL Lab. See more details at: [Link](#)

1.1 EEG paradigm


Both training and testing data are extracted from BCI Competition III - IIIb Cued motor imagery with online feedback (nonstationary classifier) with 2 classes (left hand, right hand) from 3 subjects [2 classes, 2 bipolar EEG channels] [Ref](#)

- The processed data including 2 channels, each including 1080 records and 750 timestamps.
- Take one record as example:



1.2 Depthwise Separable Convolution


This technique is proposed by [Xception](#), which is inspired by [Inception](#).



1.2.1 Low rank approximation


We could also see Depthwise Separable convolution in Linear Algebra or in Low rank approximation aspects.

Rank-1 Matrix Decomposition Considering a n by n convolution filter as a matrix, and we could apply rank-1 matrix decomposition on a matrix and get the sum of outer products of n by 1 vectors. These n by 1 vectors could be regarded as depthwise and pointwise convolution filters. Replacing the n by n convolution filters with these rank-1 vectors, we could decrease the parameters from n^2 to $2n$. Obviously, this kind of parameter reduction might decrease the capability of models.



Rank-1 Tensor Decomposition We often use high dimensional convolution filters in Computer Vision, and we need tensors for representing those high dimensional filters. Rank-1 decomposition for tensor is an optimization problem, and we could take this as extension of singular value decomposition in high dimensional domain.

rank-1 decomposition of tensor



- Practically, some deep learning frameworks like keras have implemented **separable layers** which combines depthwise convolution and pointwise convolution together, which is a depthwise layer and immediately follows by a pointwise convolution layer. We would use “separable convolution” as a term later.
- Ref: <https://medium.com/@chih.sheng.huang821/%E6%B7%B1%E5%BA%A6%E5%AD%B8%E7%BF%92mobilenet-depthwise-separable-convolution-f1ed016b3467>

2 Experimental Setup

We compare EEGNet with a Deep Convolution Network mentioned in EEGNet paper.

2.1 EEGNet architecture

- My EEGNet has the same neurons and layers with TA's EEGNet
 - Our architecture has different number of neurons with original paper, I guess TA took different some EEG data with different sample rate.

```

EEGNet(
  (firstConv): Sequential(
    (0): Conv2d(1, 16, kernel_size=(1, 51), stride=(1, 1), padding=(0, 25), bias=False)
    (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (depthwiseConv): Sequential(
    (0): Conv2d(16, 32, kernel_size=(2, 1), stride=(1, 1), groups=16, bias=False)
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU6()
    (3): AvgPool2d(kernel_size=(1, 4), stride=(1, 4), padding=0)
    (4): Dropout(p=0.2, inplace=False)
  )
  (separableConv): Sequential(
    (0): Conv2d(32, 32, kernel_size=(1, 15), stride=(1, 1), padding=(0, 7), bias=False)
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU6()
    (3): AvgPool2d(kernel_size=(1, 8), stride=(1, 8), padding=0)
    (4): Dropout(p=0.25, inplace=False)
  )
  (classify): Sequential(
    (0): Linear(in_features=736, out_features=2, bias=True)
  )
)

```

2.2 DeepConvNet architecture

- $C = 2$, $T = 750$ and $N = 2$.

```

DeepConvNet(
  (conv0): Sequential(
    (0): Conv2d(1, 25, kernel_size=(1, 5), stride=(1, 1))
    (1): Conv2d(25, 25, kernel_size=(2, 1), stride=(1, 1))
    (2): BatchNorm2d(25, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (3): ELU(alpha=1.0)
    (4): MaxPool2d(kernel_size=(1, 2), stride=(1, 2), padding=0, dilation=1, ceil_mode=False)
    (5): Dropout(p=0.5, inplace=False)
  )
  (conv1): Sequential(
    (0): Conv2d(25, 50, kernel_size=(1, 5), stride=(1, 1))
    (1): BatchNorm2d(50, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ELU(alpha=1.0)
    (3): MaxPool2d(kernel_size=(1, 2), stride=(1, 2), padding=0, dilation=1, ceil_mode=False)
    (4): Dropout(p=0.5, inplace=False)
  )
  (conv2): Sequential(
    (0): Conv2d(50, 100, kernel_size=(1, 5), stride=(1, 1))
    (1): BatchNorm2d(100, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ELU(alpha=1.0)
    (3): MaxPool2d(kernel_size=(1, 2), stride=(1, 2), padding=0, dilation=1, ceil_mode=False)
    (4): Dropout(p=0.5, inplace=False)
  )
  (conv3): Sequential(
    (0): Conv2d(100, 200, kernel_size=(1, 5), stride=(1, 1))
    (1): BatchNorm2d(200, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ELU(alpha=1.0)
    (3): MaxPool2d(kernel_size=(1, 2), stride=(1, 2), padding=0, dilation=1, ceil_mode=False)
    (4): Dropout(p=0.5, inplace=False)
  )
  (classify): Sequential(
    (0): Linear(in_features=8600, out_features=2, bias=True)
  )
)

```

3 EEGNet Details

3.1 Activation Function

3.1.1 ELU

ELU is mentioned in many papers using depthwise separable convolution like Xception, MobileNetV2, and also EEGNet. Most of these works pointed out that ReLU might much more easily suffer from the “dead neuron” problem due to smaller sizes of depthwise and pointwise layers. They recommended that it is better to use ELU activation right after pointwise convolution.

However, in my experiments, ELU performed the worst in all activation functions. I guess the reason about this phenomenon is that our dataset is so small that EEGNet does not need to learn weights in sparse space to represent the information of whole training data. Without sparsity of weights, the dead neuron problem would not be too harmful.

3.1.2 ReLU

ReLU gives good enough results.

3.1.3 ReLU6

Some empirical results found that clipping mechanism for ReLU function could make training process much more stable and get better accuracy. [Ref1](#) [Ref2](#)

3.1.4 Leaky ReLU

Not bad, but dead neurons problem didn't happen harmfully in our network so that Leaky ReLU did not outperform ReLU and ReLU6

3.2 Default Hyperparameters

- Batch size= 64
- Learning rate = 0.0025
- Epochs = 300
- Optimizer: Adam
- Loss function: `torch.nn.CrossEntropyLoss()`

4 Experimental results

4.1 Highest accuracy

- EEGNet (default dropout rate) (I tried different dropoute rates at discussion part)

	EEG_elu_train	EEG_relu_train	EEG_relu6_train	EEG_leaky_relu_train	EEG_elu_test	EEG_relu_test	EEG_relu6_test	EEG_leaky_relu_test
best_acc	99.814815	99.814815	99.537037	99.722222	87.12963	87.12963	87.314815	88.240741

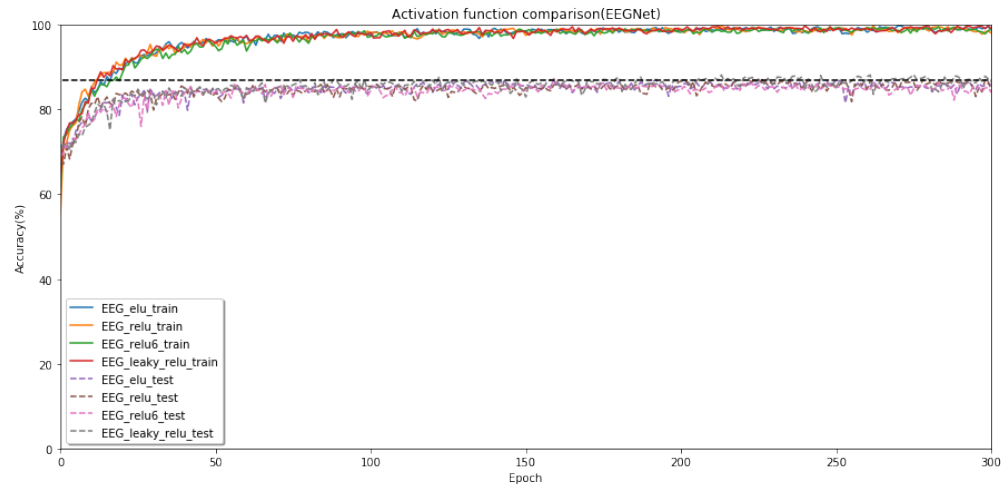
- DeepConvNet

	DCN_elu_train	DCN_relu_train	DCN_relu6_train	DCN_leaky_relu_train	DCN_elu_test	DCN_relu_test	DCN_relu6_test	DCN_leaky_relu_test
best_acc	98.055556	97.037037	96.851852	97.222222	81.851852	80.833333	81.111111	82.777778

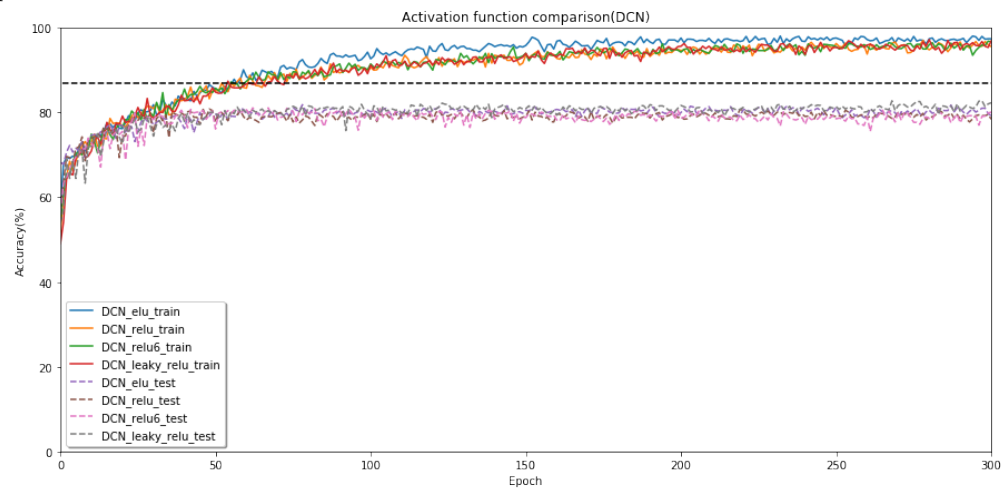
4.2 Result Comparison

There is a black dotted line at 87%

- EEGNet



- DeepConvNet



5 Discussion and extra experiments

5.1 Hyperparameter Tuning

Hyperparameter Tuning is only applied to EEGNet because EEGNet is much more practical than the DeepConvNet, and there are less parameters in EEGNet so that this tuning process could be faster.

- Though dropout rates don't belong to hyperparameters, I tried to tune dropout rates together to see whether it could improve accuracy

5.1.1 Ray Tune

[Ray Tune](#) is a state-of-the-art distributed hyperparameter tuning framework, supporting many famous tuning algorithms including grid search, [Population Based Training \(PBT\)](#), and also supporting some early stopping algorithms like [Asynchronous Successive Halving Algorithm \(ASHA\)](#)

- Using Early stopping algorithm often stops too early to get better accuracy, so I omit this.

```
== Status ==
Memory usage on this node: 31.0/31.0 GiB: ***LOW MEMORY*** less than 10% of the memory on this node is available for use. This can cause unexpected crashes. Consider reducing the memory used by your application or reducing the Ray object store size by setting 'object_store_memory' when calling 'ray.init'.
Using AsyncHyperBand: num_stopped=28
Bracket: Iter 64.000: 0.762962962962963 | Iter 16.000: 0.7171296296296297 | Iter 4.000: 0.712037037037037 | Iter 1.000: 0.6972222222222222
Resources requested: 3/12 CPUs, 1/2 GPUs, 0.0/5.62 GiB heap, 0.0/1.9 GiB objects
Result logdir: /home/allurus/ray_results/TrainEEG
Number of trials: 120 (108 ERROR, 1 RUNNING, 11 TERMINATED)
```

Trial name	status	loc	drop1	drop2	lr	acc	iter	total time (s)
TrainEEG_00009	ERROR		0.65	0.65	0.003	0.707407	1	1.44206
TrainEEG_00010	ERROR		0.6	0.68	0.003			
TrainEEG_00011	ERROR		0.65	0.68	0.003	0.697222	1	1.3442
TrainEEG_00012	ERROR		0.6	0.65	0.0022			
TrainEEG_00013	ERROR		0.65	0.65	0.0022			
TrainEEG_00014	ERROR		0.6	0.68	0.0022			
TrainEEG_00015	ERROR		0.65	0.68	0.0022			
TrainEEG_00016	ERROR		0.6	0.65	0.0025			
TrainEEG_00017	ERROR		0.65	0.65	0.0025			
TrainEEG_00018	ERROR		0.6	0.68	0.0025			
TrainEEG_00119	RUNNING		0.65	0.68	0.003			
TrainEEG_00000	TERMINATED		0.6	0.65	0.0022	0.65463	1	1.69158
TrainEEG_00001	TERMINATED		0.65	0.65	0.0022	0.781481	100	6.93045
TrainEEG_00002	TERMINATED		0.6	0.68	0.0022	0.700926	4	1.58215
TrainEEG_00003	TERMINATED		0.65	0.68	0.0022	0.681481	1	1.57006
TrainEEG_00004	TERMINATED		0.6	0.65	0.0025	0.664815	1	1.52389
TrainEEG_00005	TERMINATED		0.65	0.65	0.0025	0.706481	4	1.58193
TrainEEG_00006	TERMINATED		0.6	0.68	0.0025	0.667593	1	1.60679
TrainEEG_00007	TERMINATED		0.65	0.68	0.0025	0.700926	4	1.70325
TrainEEG_00008	TERMINATED		0.6	0.65	0.003	0.691667	4	1.73266
TrainEEG_00046	TERMINATED		0.6	0.68	0.003	0.686111	1	1.29691

- Ray Tune running Example

```
== Status ==
Memory usage on this node: 4.5/88.5 GiB
Using FIFO scheduling algorithm.
Resources requested: 3/4 CPUs, 1/1 GPUs, 0.0/52.83 GiB heap, 0.0/18.21 GiB objects
Result logdir: /home/ubuntu/ray_results/TrainEEG
Number of trials: 800 (1 RUNNING, 799 TERMINATED)
```

Trial name	status	loc	activation	drop1	drop2	lr	acc	iter	total time (s)
TrainEEG_00799	RUNNING	10.54.87.6:3614	3	0.65	0.68	0.0039	0.836111	256	12.6855
TrainEEG_00000	TERMINATED		0	0.6	0.65	0.0022	0.815741	280	13.4511
TrainEEG_00001	TERMINATED		1	0.6	0.65	0.0022	0.860185	280	13.4434
TrainEEG_00002	TERMINATED		2	0.6	0.65	0.0022	0.873148	85	5.22564
TrainEEG_00003	TERMINATED		3	0.6	0.65	0.0022	0.855556	280	14.1445
TrainEEG_00004	TERMINATED		0	0.65	0.65	0.0022	0.828704	280	13.9425
TrainEEG_00005	TERMINATED		1	0.65	0.65	0.0022	0.872222	260	13.3257
TrainEEG_00006	TERMINATED		2	0.65	0.65	0.0022	0.872222	209	10.7938
TrainEEG_00007	TERMINATED		3	0.65	0.65	0.0022	0.844444	280	13.7365
TrainEEG_00008	TERMINATED		0	0.6	0.68	0.0022	0.799074	280	13.9066
TrainEEG_00009	TERMINATED		1	0.6	0.68	0.0022	0.862963	280	14.7089
TrainEEG_00010	TERMINATED		2	0.6	0.68	0.0022	0.846296	280	13.725
TrainEEG_00011	TERMINATED		3	0.6	0.68	0.0022	0.843519	280	13.4502
TrainEEG_00012	TERMINATED		0	0.65	0.68	0.0022	0.84537	280	13.8676
TrainEEG_00013	TERMINATED		1	0.65	0.68	0.0022	0.871296	264	12.878
TrainEEG_00014	TERMINATED		2	0.65	0.68	0.0022	0.856481	280	13.6226
TrainEEG_00015	TERMINATED		3	0.65	0.68	0.0022	0.851852	280	14.0215
TrainEEG_00016	TERMINATED		0	0.6	0.65	0.0025	0.814815	280	13.5289
TrainEEG_00017	TERMINATED		1	0.6	0.65	0.0025	0.87037	174	8.819
TrainEEG_00018	TERMINATED		2	0.6	0.65	0.0025	0.859259	280	13.3686

... 780 more trials not shown (780 TERMINATED)

- Ray Tune result Example

TrainEEG_00770	TERMINATED		2	0.6	0.65	0.0035	0.857407	280	13.8504
TrainEEG_00771	TERMINATED		3	0.6	0.65	0.0035	0.842593	280	13.5983
TrainEEG_00772	TERMINATED		0	0.65	0.65	0.0035	0.806481	280	13.6223
TrainEEG_00773	TERMINATED		1	0.65	0.65	0.0035	0.857407	280	13.6689
TrainEEG_00774	TERMINATED		2	0.65	0.65	0.0035	0.873148	174	9.20229
TrainEEG_00775	TERMINATED		3	0.65	0.65	0.0035	0.833333	280	13.327
TrainEEG_00776	TERMINATED		0	0.6	0.68	0.0035	0.824074	280	13.782
TrainEEG_00777	TERMINATED		1	0.6	0.68	0.0035	0.85463	280	13.4184
TrainEEG_00778	TERMINATED		2	0.6	0.68	0.0035	0.857407	280	13.4881
TrainEEG_00779	TERMINATED		3	0.6	0.68	0.0035	0.852778	280	13.1828
TrainEEG_00780	TERMINATED		0	0.65	0.68	0.0035	0.784259	280	13.4409
TrainEEG_00781	TERMINATED		1	0.65	0.68	0.0035	0.85	280	13.4271
TrainEEG_00782	TERMINATED		2	0.65	0.68	0.0035	0.853704	280	14.3059
TrainEEG_00783	TERMINATED		3	0.65	0.68	0.0035	0.863889	280	13.6275
TrainEEG_00784	TERMINATED		0	0.6	0.65	0.0039	0.807407	280	14.0187
TrainEEG_00785	TERMINATED		1	0.6	0.65	0.0039	0.853704	280	13.3928
TrainEEG_00786	TERMINATED		2	0.6	0.65	0.0039	0.844444	280	13.2157
TrainEEG_00787	TERMINATED		3	0.6	0.65	0.0039	0.864815	280	13.6354
TrainEEG_00788	TERMINATED		0	0.65	0.65	0.0039	0.799074	280	13.3651
TrainEEG_00789	TERMINATED		1	0.65	0.65	0.0039	0.828704	280	13.4807
TrainEEG_00790	TERMINATED		2	0.65	0.65	0.0039	0.863889	280	13.7216
TrainEEG_00791	TERMINATED		3	0.65	0.65	0.0039	0.872222	164	8.50717
TrainEEG_00792	TERMINATED		0	0.6	0.68	0.0039	0.816667	280	13.4541
TrainEEG_00793	TERMINATED		1	0.6	0.68	0.0039	0.87037	209	10.4741
TrainEEG_00794	TERMINATED		2	0.6	0.68	0.0039	0.873148	227	11.1149
TrainEEG_00795	TERMINATED		3	0.6	0.68	0.0039	0.871296	204	10.102
TrainEEG_00796	TERMINATED		0	0.65	0.68	0.0039	0.816667	280	13.4023
TrainEEG_00797	TERMINATED		1	0.65	0.68	0.0039	0.848148	280	13.5914
TrainEEG_00798	TERMINATED		2	0.65	0.68	0.0039	0.866667	280	13.3319
TrainEEG_00799	TERMINATED		3	0.65	0.68	0.0039	0.84537	280	13.8336

Best config is: {'activation': 1, 'args': Namespace(ray_address=None, smoke_test=False, use_gpu=True), 'drop1': 0.6, 'drop2': 0.65, 'lr': 0.0025}

5.2 Tuning Results

5.2.1 Adam with weight decay

Due to the learning rate is set to be 0.0028, weight decay should be less than learning rate to improve optimization process.

```
1 colname = ['lr=.0028', 'lr=.0028, wd=.01', 'lr=.0028, wd=.0001']
2 max_list = [df8.max(), df8_01.max(), df8_0001.max()]
3 pd.DataFrame(max_list, index=colname).T.iloc[4:]
```

	lr=.0028	lr=.0028, wd=.01	lr=.0028, wd=.0001
EEG_elu_test	83.796296	71.944444	85.555556
EEG_relu_test	87.962963	81.481481	89.259259
EEG_relu6_test	88.055556	73.888889	88.333333
EEG_leaky_relu_test	87.777778	72.314815	88.425926

5.2.2 Best hyperparameters

By Ray Tune, we get best learning rate between [0.0022 to 0.0028]. I apply the weight decay rate to be 0.001 to see whether it could be much better.

Surprisingly, ELU performs usually bad, while ReLU6 often outperform ReLU.

```
1 import pandas as pd
2 df2 = pd.read_csv('eeg_0022.csv', index_col=0)
3 df5 = pd.read_csv('eeg_0025.csv', index_col=0)
4 df8 = pd.read_csv('eeg_0028.csv', index_col=0)
5 df8_01 = pd.read_csv('eeg_0028_wd_01.csv', index_col=0)
6 df8_0001 = pd.read_csv('eeg_0028_wd_0001.csv', index_col=0)
```

```
1 colname = ['lr=.0022', 'lr=.0025', 'lr=.0028']
2 max_list = [df2.max(), df5.max(), df8.max()]
3 pd.DataFrame(max_list, index=colname).T.iloc[4:]
```

	lr=.0022	lr=.0025	lr=.0028
EEG_elu_test	85.555556	85.185185	83.796296
EEG_relu_test	86.759259	88.148148	87.962963
EEG_relu6_test	88.703704	89.259259	88.055556
EEG_leaky_relu_test	88.240741	87.685185	87.777778

```
1 colname = ['lr=.0025', 'lr=.0025, wd=.0001']
2 max_list = [df5.max(), df5_0001.max()]
3 pd.DataFrame(max_list, index=colname).T.iloc[4:]
```

	lr=.0025	lr=.0025, wd=.0001
EEG_elu_test	85.185185	84.259259
EEG_relu_test	88.148148	86.574074
EEG_relu6_test	<u>89.259259</u>	88.611111
EEG_leaky_relu_test	87.685185	<u>89.166667</u>

5.3 Conclusion

EEGNet is an excellent network with efficient training time and good enough performance. Actually, the first version of EEGNet posted on arxiv at 2016 does not include Depthwise Separable Convolution. It is until the version3 and version4 at 2018 that EEGNet take some experiences from MobileNet and Xception so that this work become better. In EEG classification problem, or multi-channel time series classification, EEGNet is still state-of-the-art. However, I think EEGNet could be extended more, but maybe there are less researcher in BCI discipline working with it. Maybe some related works of Stock Market prediction, or some tricks in MobileNetV3 could apply on EEGNet to improve it. I would try this in my final project.

6 Code

- For recording accuracy, we provide a train function and a test function in main.py
- In tuning process, it's not necessary to record accuracy of each epoch so that we provide simpler train function and test function in EEGNet.py

6.1 main.py

```
[ ]: from functools import reduce

import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.optim
import pandas as pd
from torch.cuda import device
from torch.utils.data import TensorDataset, DataLoader

from dataloader import read_bci_data
from EEGNet import EEGNet
from DCN2 import DeepConvNet

def get_bci_dataloaders():
    train_x, train_y, test_x, test_y = read_bci_data()
    datasets = []
    for train, test in [(train_x, train_y), (test_x, test_y)]:
        train = torch.stack(
            [torch.Tensor(train[i]) for i in range(train.shape[0])]
        )
        test = torch.stack(
            [torch.Tensor(test[i:i+1]) for i in range(test.shape[0])]
        )
        datasets += [TensorDataset(train, test)]

    return datasets

def get_data_loaders(train_dataset, test_dataset):
    #train_dataset, test_dataset = get_bci_dataloaders()
    kwargs = {'num_workers': 1, 'pin_memory': True} if torch.cuda.is_available()
    else {}
    train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
    test_loader = DataLoader(test_dataset, batch_size=len(test_dataset),
    shuffle=True)
    return train_loader, test_loader

def showResult(title='', **kwargs):
```

```

plt.figure(figsize=(15, 7))
plt.title(title)
plt.xlabel('Epoch')
plt.ylabel('Accuracy(%)')

for label, data in kwargs.items():
    plt.plot(range(len(data)), data, '--' if 'test' in label else '-',
→label=label)
    plt.ylim(0, 100)
    plt.xlim(0, 300)
    points = [(-5, 87), (310, 87)]
    (xpoints, ypoints) = zip(*points)

plt.plot(xpoints, ypoints, linestyle='--', color='black')

plt.legend(loc='best', fancybox=True, shadow=True)
plt.show()

def main():
    torch.backends.cudnn.enabled = True
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

    nets1 = {
        "EEG_elu": EEGNet(nn.ELU).to(device),
        "EEG_relu": EEGNet(nn.ReLU).to(device),
        "EEG_relu6": EEGNet(nn.ReLU6).to(device),
        "EEG_leaky_relu": EEGNet(nn.LeakyReLU).to(device)
    }

    nets2 = {
        "DCN_elu": DeepConvNet(nn.ELU).to(device),
        "DCN_relu": DeepConvNet(nn.ReLU).to(device),
        "DCN_relu6": DeepConvNet(nn.ReLU6).to(device),
        "DCN_leaky_relu": DeepConvNet(nn.LeakyReLU).to(device)
    }

    #nets = nets1
    nets = nets2

    # Training setting
    loss_fn = nn.CrossEntropyLoss()
    learning_rates = {0.0025}

    optimizer = torch.optim.Adam
    optimizers = {

```

```

        key: optimizer(value.parameters(), lr=learning_rate, weight_decay=0.0001)
    for key, value in nets.items()
    for learning_rate in learning_rates
}

epoch_size = 300
batch_size = 64
acc = train(nets, epoch_size, batch_size, loss_fn, optimizers)
df = pd.DataFrame.from_dict(acc)

#df.to_csv('eeg_0025_0001.csv')
#print(df)
display(df)
return df

# This train is for demo and recording accuracy
def train(nets, epoch_size, batch_size, loss_fn, optimizers):
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

    trainDataset, testDataset = get_bci_dataloaders()
    trainLoader, testLoader = get_data_loaders(trainDataset, testDataset)

    accuracy = {
        **{key + "_train": [] for key in nets},
        **{key + "_test": [] for key in nets}
    }
    for epoch in range(epoch_size + 1):
        train_correct = {key: 0.0 for key in nets}
        test_correct = {key: 0.0 for key in nets}
        for step, (x, y) in enumerate(trainLoader):
            x = x.to(device)
            y = y.to(device).long().view(-1)

            for key, net in nets.items():
                net.train(mode=True)
                y_hat = net(x)
                loss = loss_fn(y_hat, y)
                loss.backward()
                train_correct[key] += (torch.max(y_hat, 1)[1] == y).sum().item()

            for optimizer in optimizers.values():
                optimizer.step()
                optimizer.zero_grad()

        with torch.no_grad():
            for step, (x, y) in enumerate(testLoader):
                x = x.to(device)

```

```

        y = y.to(device).long().view(-1)
        for key, net in nets.items():
            net.eval()
            y_hat = net(x)
            test_correct[key] += (torch.max(y_hat, 1)[1] == y).sum().
→ item()

    for key, value in train_correct.items():
        accuracy[key + "_train"] += [(value * 100.0) / len(trainDataset)]

    for key, value in test_correct.items():
        accuracy[key + "_test"] += [(value * 100.0) / len(testDataset)]

    if epoch % 100 == 0:
        print('epoch : ', epoch, ' loss : ', loss.item())
        display(pd.DataFrame.from_dict(accuracy).iloc[[epoch]])
        print('')
        torch.cuda.empty_cache()
        showResult(title='Activation function comparison(EEGNet)'.format(epoch + 1),
→ **accuracy)
        #showResult(title='Activation function comparison(DCN)'.format(epoch + 1),
→ **accuracy)
    return accuracy

if __name__ == '__main__':
    df1 = main()
    df1.max()
    display(pd.DataFrame(df1.max(), columns=['best_acc']).T)

```

6.2 EEGNet.py

```

[ ]: import torch
import torch.nn as nn
import torch.optim
import torch.nn.functional as F
from torch.cuda import device
from torch.utils.data import DataLoader

EPOCH_SIZE = 512
TEST_SIZE = 256

class EEGNet(nn.Module):
    def __init__(self, activation=None, dropout1=0.25, dropout2=0.25):
        if not activation:
            activation = nn.ELU
        super(EEGNet, self).__init__()

```

```

        self.firstConv = nn.Sequential(
            nn.Conv2d(1, 16, kernel_size=(1, 51), stride=(1, 1), padding=(0,
→25), bias=False),
            nn.BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
→track_running_stats=True)
        )

        self.depthwiseConv = nn.Sequential(
            nn.Conv2d(16, 32, kernel_size=(2, 1), stride=(1, 1), groups=16,
→bias=False),
            nn.BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
→track_running_stats=True),
            activation(),
            nn.AvgPool2d(kernel_size=(1, 4), stride=(1, 4), padding=0),
            nn.Dropout(p=dropout1)
        )

        self.separableConv = nn.Sequential(
            nn.Conv2d(32, 32, kernel_size=(1, 15), stride=(1, 1), padding=(0,
→7), bias=False),
            nn.BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
→track_running_stats=True),
            activation(),
            nn.AvgPool2d(kernel_size=(1, 8), stride=(1, 8), padding=0),
            nn.Dropout(p=dropout2)
        )

        self.classify = nn.Sequential(
            nn.Linear(in_features=736, out_features=2, bias=True)
        )

    def forward(self, x):
        x = self.firstConv(x)
        x = self.depthwiseConv(x)
        x = self.separableConv(x)
        x = x.view(-1, self.classify[0].in_features)
        x = self.classify(x)
        return x

# This train func is for tuning only so that accuracy recoding is removed
def train(model, optimizer, train_loader, device=torch.device("cpu")):
    model.train()
    loss_fn = nn.CrossEntropyLoss()
    for batch_idx, (data, target) in enumerate(train_loader):
        if batch_idx * len(data) > EPOCH_SIZE:

```

```

        return
    data, target = data.to(device), target.to(device)
    optimizer.zero_grad()
    output = model(data)
    loss = loss_fn(output, target.squeeze().long())
    loss.backward()
    optimizer.step()

# This test func is for tuning only so that accuracy recoding is removed
def test(model, test_loader, device=torch.device("cpu")):
    model.eval()
    test_loss = 0
    correct = 0
    total = 0
    with torch.no_grad():
        for batch_idx, (data, target) in enumerate(test_loader):
            if batch_idx * len(data) > TEST_SIZE:
                break
            data, target = data.to(device), target.squeeze().to(device)
            outputs = model(data)
            _, predicted = torch.max(outputs.data, 1)
            total += target.size(0)
            correct += (predicted == target).sum().item()

    return correct / total

```

6.3 DCN.py

```

[ ]: from functools import reduce
import torch
import torch.nn as nn
import torch.optim as optim

class DeepConvNet(nn.Module):
    def __init__(self, activation=None, deepconv=[25,50,100,200], dropout=0.5):
        super(DeepConvNet, self).__init__()

        if not activation:
            activation = nn.ELU

        self.deepconv = deepconv
        self.conv0 = nn.Sequential(
            nn.Conv2d(
                1, deepconv[0], kernel_size=(1, 5),
                stride=(1,1), padding=(0,0), bias=True
            ),

```



```

        nn.Conv2d(
            deepconv[0], deepconv[0], kernel_size=(2,1),
            stride=(1,1), padding=(0,0), bias=True
        ),
        nn.BatchNorm2d(deepconv[0]),
        activation(),
        nn.MaxPool2d(kernel_size=(1,2)),
        nn.Dropout(p=dropout)
    )

    for idx in range(1, len(deepconv)):
        setattr(self, 'conv'+str(idx), nn.Sequential(
            nn.Conv2d(
                deepconv[idx-1], deepconv[idx], kernel_size=(1,5),
                stride=(1,1), padding=(0,0), bias=True
            ),
            nn.BatchNorm2d(deepconv[idx]),
            activation(),
            nn.MaxPool2d(kernel_size=(1, 2)),
            nn.Dropout(p=dropout)
        ))

    flatten_size = deepconv[-1] * reduce(
        lambda x, _: round((x-4)/2), deepconv, 750)
    self.classify = nn.Sequential(
        nn.Linear(flatten_size, 2, bias=True),
    )

    def forward(self, x):
        for i in range(len(self.deepconv)):
            x = getattr(self, 'conv'+str(i))(x)
        # flatten
        x = x.view(-1, self.classify[0].in_features)
        x = self.classify(x)
        return x

```

6.4 Tuning.py

This code should run in command line interface to get arguments

```
[ ]:
```

```
[ ]: from __future__ import print_function
```

```
import argparse
import os
```

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
from dataloader import read_bci_data
from EEGNet import (EEGNet, train, test)

import ray
from ray import tune
from ray.tune.schedulers import ASHAScheduler

activation_list = [nn.ELU, nn.ReLU, nn.ReLU6, nn.LeakyReLU]

def gen_bci_dataloaders():
    train_x, train_y, test_x, test_y = read_bci_data()
    datasets = []
    for train, test in [(train_x, train_y), (test_x, test_y)]:
        train = torch.stack(
            [torch.Tensor(train[i]) for i in range(train.shape[0])]
        )
        test = torch.stack(
            [torch.Tensor(test[i:i+1]) for i in range(test.shape[0])]
        )
        datasets += [TensorDataset(train, test)]

    return datasets

def get_data_loaders():
    train_dataset, test_dataset = gen_bci_dataloaders()
    kwargs = {'num_workers': 1, 'pin_memory': True} if torch.cuda.is_available()
    else {}
    train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
    test_loader = DataLoader(test_dataset, batch_size=len(test_dataset),
    shuffle=True)
    return train_loader, test_loader

# Change these values if you want the training to run quicker or slower.
EPOCH_SIZE = 512
TEST_SIZE = 256

# Training settings
parser = argparse.ArgumentParser(description="PyTorch MNIST Example")
parser.add_argument(
    "--use-gpu", action="store_true", default=True,
    help="enables CUDA training")
parser.add_argument(
    "--ray-address", type=str, help="The Redis address of the cluster.")

```

```

parser.add_argument(
    "--smoke-test", action="store_true", help="Finish quickly for testing")

class TrainEEG(tune.Trainable):
    def _setup(self, config):
        self.device = torch.device("cuda" if torch.cuda.is_available() else
→"cpu")
        self.train_loader, self.test_loader = get_data_loaders()
        self.model = EEGNet(
            activation=activation_list[config.get("activation", 1)]).to(self.
→device)

            #dropout1=config.get("drop1", 0.6),
            #dropout2=config.get("drop2", 0.65)
            #).to(self.device)

        self.model = self.model.to(self.device)
        self.optimizer = optim.Adam(
            self.model.parameters(),
            lr=config.get("lr", 0.0025))

    def _train(self):
        train(
            self.model, self.optimizer, self.train_loader, device=self.device)
        acc = test(self.model, self.test_loader, self.device)
        return {"mean_accuracy": acc}

    def _save(self, checkpoint_dir):
        checkpoint_path = os.path.join(checkpoint_dir, "model.pth")
        torch.save(self.model.state_dict(), checkpoint_path)
        return checkpoint_path

    def _restore(self, checkpoint_path):
        self.model.load_state_dict(torch.load(checkpoint_path))

if __name__ == "__main__":
    return
    args = parser.parse_args()
    ray.init(address=args.ray_address, num_cpus=3 if args.smoke_test else None)
    #sched = ASHAScheduler(metric="mean_accuracy")

    analysis = tune.run(
        TrainEEG,
        #scheduler=sched,

```

```

stop={
    "mean_accuracy": 0.87,
    "training_iteration": 3 if args.smoke_test else 300,
},
resources_per_trial={
    "cpu": 3,
    "gpu": int(args.use_gpu)
},
num_samples=1 if args.smoke_test else 10,
checkpoint_at_end=True,
checkpoint_freq=10,
config={
    "activation": tune.grid_search([0, 1, 2, 3]),
    # "args": args,
    # "drop1": tune.grid_search([0.6, 0.65]),
    # "drop2": tune.grid_search([0.65, 0.68]),
    "lr": tune.grid_search([0.0025, 0.0026])
    # "lr": tune.grid_search([0.0022, 0.0025, 0.003, 0.01])
    # "lr": tune.uniform(0.0022, 0.005)
})

print("Best config is:", analysis.get_best_config(metric="mean_accuracy"))

```